

# Cifar10-CNN ipynb version

Image Classification using Convolutional Neural Networks(CNN) in PyTorch

## 1. Exploring the Cifar10 Dataset

```
In [1]: import os
import torch
import torchvision
import tarfile
from torchvision.datasets.utils import download_url
from torch.utils.data import random_split

#UsersInfor>rmnet>amazon>aws>wget>url>lib>site>packages>torch>auto.py:22: TqdmWarning: IPProgress not from
C: Please update jupyter and ipynb kernels. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
/UsersInfor>rmnet>amazon>aws>wget>url>lib>site>packages>torch>vision>io>image.py:13: UserWarning: Failed
to load image Python extension:
warn(Failed to load image Python extension: (e*))

In [2]: # Download the dataset
dataset_url = "https://s3.amazonaws.com/fast-ai-imageclas/cifar10.tgz"
download_url(dataset_url, '.')

Using downloaded and verified file: .\cifar10.tgz

In [3]: # Extract from archive
with tarfile.open('.\cifar10.tgz', 'r:gz') as tar:
    tar.extractall(path='.')

In [4]: data_dir = './data/cifar10'

import torchvision
print('Listdir',data_dir)
classes = os.listdir(data_dir + '/train')
print(classes)

['test', 'train']
['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

두개의 폴더 내부폴더 살펴보면, 하나는 train, test이고 각 클래스에 들어있는 이미지 수가 있는지 확인. train은 5000개, test는 1000개

In [5]: airplane_files = os.listdir(data_dir + '/train/airplane')
print('No. of training examples for airplane:', len(airplane_files))
print(airplane_files[:5])

No. of training examples for airplanes: 5000
['0001.png', '0002.png', '0003.png', '0004.png', '0005.png']

In [6]: ship_test_files = os.listdir(data_dir + '/test/ship')
print('No. of test examples for ship:', len(ship_test_files))
print(ship_test_files[:5])

No. of test examples for ship: 1000
['0001.png', '0002.png', '0003.png', '0004.png', '0005.png']

이미지 폴더의 class를 torchvision을 통해 pytorch tensor로 로드합니다.

In [7]: from torchvision.datasets import ImageFolder
from torchvision.transforms import transforms

In [8]: dataset = ImageFolder(data_dir+'/'+'train', transform = ToTensor())

각 클래스는 폴더이름이 이미지 train 폴더와 동일합니다.
airplane, ship, cat, horse, frog, deer, dog, truck, bird, automobile
이 이미지 텐서는 32x32 pixel에 channel은 3(RGB)과 size가 100 이미지 shape = (3,32,32)
```

```
In [9]: img, label = dataset[0]
print(img.size(), label)

torch.Size([3, 32, 32]) 0
tensor([[[[ 0.7922, 0.7922, 0.8000, ..., 0.8118, 0.8039, 0.7765]],
[ 0.8078, 0.8075, 0.8318, ..., 0.8235, 0.8357, 0.7891]],
[ 0.8225, 0.8275, 0.8314, ..., 0.8392, 0.8314, 0.8235]],
...,
[ 0.8549, 0.8235, 0.7668, ..., 0.9529, 0.9569, 0.9529]],
[ 0.8588, 0.8510, 0.8471, ..., 0.9451, 0.9451, 0.9451]],
[ 0.8510, 0.8471, 0.8518, ..., 0.9373, 0.9373, 0.9412]],
...,
[ 0.8900, 0.8990, 0.8918, ..., 0.8157, 0.8078, 0.8000]],
[ 0.8157, 0.8157, 0.8196, ..., 0.8275, 0.8196, 0.8118]],
[ 0.8314, 0.8353, 0.8392, ..., 0.8392, 0.8392, 0.8275]],
...,
[ 0.7894, 0.7884, 0.7882, ..., 0.7843, 0.7834, 0.7765]],
[ 0.7951, 0.7964, 0.8000, ..., 0.8032, 0.7954, 0.7891]],
[ 0.8118, 0.8161, 0.8235, ..., 0.8355, 0.8357, 0.8078]],
...,
[ 0.7760, 0.8392, 0.7795, ..., 0.6995, 0.6985, 0.6985]],
[ 0.8745, 0.8857, 0.8627, ..., 0.9688, 0.9688, 0.9687]],
[ 0.8667, 0.8627, 0.8667, ..., 0.9529, 0.9529, 0.9529]])]

In [10]: print(dataset.classes)

['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

이 pytorch 텐서에서 시각화를 하고 싶다면 matplotlib을 사용하여 데 이터는 텐서 차원을
(32x32x3)으로 바꿔야함

In [11]: import matplotlib.pyplot as plt
import matplotlib inline

fig=plt.figure(figsize=(10,10))
plt.imshow(img.permute(1,2,0))

plt.colorbar()

plt.title('airplane')

In [12]: def show_example(img, label):
    print('Label:', dataset.classes[label], '(*' + str(label) + ')')
    plt.imshow(img.permute(1,2,0))

    # Permute() : 모든 차원들을 뒤로함할 수 있음
    # ex)
    # x = torch.rand(3, 32, 3)
    # x = x.permute(1,2,0)

In [13]: show_example(dataset[0])

Label: airplane (0)

0
5
10
15
20
25
30
0 5 10 15 20 25 30

In [14]: show_example(dataset[1000])

Label: bird (2)

0
5
10
15
20
25
30
0 5 10 15 20 25 30
```

## Train, Valid Split

```
In [15]: random_seed = 42
torch.manual_seed(random_seed);

In [16]: val_size = 5000
train_ds, val_ds = random_split(dataset, [train_size, val_size])
len(train_ds), len(val_ds)

Out[16]: (4500, 5000)

In [17]: from torch.utils.data.dataloader import DataLoader

batch_size=256

배치사이즈에 load 시켜주려면 data loader가 필요하기까

In [18]: train_dl = DataLoader(train_ds, batch_size, shuffle=True, num_workers=4, pin_memory=True)
val_dl = DataLoader(val_ds, batch_size=2, num_workers=4, pin_memory=True)

In [19]: from torchvision.utils import make_grid

def show_batch(dl):
    for images, labels in dl:
        fig, ax = plt.subplots(figsize=(12,6))
        ax.set_xticks([]); ax.set_yticks([])
        ax.imshow(make_grid(images, nrow=6).permute(1, 2, 0))
        break

In [20]: show_batch(train_dl)


```

GPU를 활용하게 사용하기 위해 사용 가능한 경우 두 가지 도우미 함수(torch\_default\_device 및 to\_device)와 도우미 클래스
DeviceDataLoader를 정의하여 필요에 따라 모델 및 데이터로 GPU로 이동시킨다

```
In [21]: def get_default_device():
    """Pick GPU if available, else CPU"""
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    """Move tensor(s) to chosen device"""
    if isinstance(data, (list,tuple)):
        return [to_device(d, device) for x in data]
    return data.to(device, device != torch.device('cpu'))

class DeviceDataLoader():
    """Wrap a dataloader to move data to a device"""
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        """Yield a batch of data after moving it to device"""
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        """Number of batches"""
        return len(self.dl)

In [30]: device = get_default_device()
device

Out[30]: device(type='cuda')
```

이제 DeviceDataLoader를 사용하여 데이터로더를 GPU로 자동 전송(사용 가능한 경우)하고 to\_device를 사용하여 모델을 GPU(사
용 가능한 경우)로 이동하기 위해 교착 및 검증 데이터 로더를 래핑할 수 있습니다.

```
In [31]: train_loader = DeviceDataLoader(train_dl, device)
val_loader = DeviceDataLoader(val_dl, device)
to_device(model, device);

Training the Model

In [32]: @torch.no_grad()
def evaluate(model, val_loader):
    model.eval()
    num_batches = len(val_loader)
    train_loader.validation_step(batch) for batch in val_loader]
    return model.validation_epoch_end(outputs)

def fit(epochs, lr, model, train_loader, val_loader, opt_func=torch.optim.SGD):
    history = []
    optimizer = opt_func(model.parameters(), lr)
    for epoch in range(epochs):
        # Training Phase
        model.train()
        train_losses = []
        for batch in train_loader:
            loss = model.training_step(batch)
            train_losses.append(loss)
        loss.backward()
        optimizer.step()
        optimizer.zero_grad()
        # Validation Phase
        result = evaluate(model, val_loader)
        result['train_loss'] = torch.stack(train_losses).mean().item()
        model.validation_step(batch) result
        history.append(result)
    return history

In [33]: model = to_device(Cifar10CNNModel(), device)

In [34]: evaluate(model, val_loader)

Out[34]: ('val_loss': 2.3023259102996826, 'val_acc': 0.69949378669261932)
```

무작위적인 샘플링(초기 정확도는 10% 이하)을 위해 모델은 훈련하기 위해 하위퍼파라미터(학습속도, epoch수, batch\_size 등을)
사용할 것이다.> 높은 정확도를 위하여

```
In [35]: num_epochs = 200
opt_func = torch.optim.AdamGrad
lr = 0.001

In [36]: history = fit(num_epochs, lr, model, train_dl, val_dl, opt_func)

Epoch [0], train_loss: 1.9730, val_loss: 1.7252, val_acc: 0.3597
Epoch [1], train_loss: 1.6415, val_loss: 1.5660, val_acc: 0.4234
Epoch [2], train_loss: 1.5394, val_loss: 1.5609, val_acc: 0.4438
Epoch [3], train_loss: 1.4647, val_loss: 1.4388, val_acc: 0.4764
Epoch [4], train_loss: 1.4160, val_loss: 1.4121, val_acc: 0.4876
Epoch [5], train_loss: 1.3730, val_loss: 1.3604, val_acc: 0.5090
Epoch [6], train_loss: 1.3363, val_loss: 1.3475, val_acc: 0.5193
Epoch [7], train_loss: 1.3025, val_loss: 1.2974, val_acc: 0.5313
Epoch [8], train_loss: 1.2737, val_loss: 1.2781, val_acc: 0.5438
Epoch [9], train_loss: 1.2467, val_loss: 1.2480, val_acc: 0.5535
Epoch [10], train_loss: 1.2201, val_loss: 1.2354, val_acc: 0.5585
Epoch [11], train_loss: 1.1921, val_loss: 1.2151, val_acc: 0.5659
Epoch [12], train_loss: 1.1821, val_loss: 1.1961, val_acc: 0.5721
Epoch [13], train_loss: 1.1627, val_loss: 1.1939, val_acc: 0.5793
Epoch [14], train_loss: 1.1453, val_loss: 1.1791, val_acc: 0.5859
Epoch [15], train_loss: 1.1291, val_loss: 1.1561, val_acc: 0.5871
Epoch [16], train_loss: 1.1132, val_loss: 1.1308, val_acc: 0.5881
Epoch [17], train_loss: 1.0997, val_loss: 1.1094, val_acc: 0.5892
Epoch [18], train_loss: 1.0867, val_loss: 1.1476, val_acc: 0.5883
Epoch [19], train_loss: 1.0693, val_loss: 1.1188, val_acc: 0.5963
Epoch [20], train_loss: 1.0559, val_loss: 1.1059, val_acc: 0.6058
Epoch [21], train_loss: 1.0426, val_loss: 1.0953, val_acc: 0.6086
Epoch [22], train_loss: 1.0329, val_loss: 1.0988, val_acc: 0.6132
Epoch [23], train_loss: 1.0196, val_loss: 1.0830, val_acc: 0.6192
Epoch [24], train_loss: 1.0066, val_loss: 1.0789, val_acc: 0.6208
Epoch [25], train_loss: 0.9986, val_loss: 1.0631, val_acc: 0.6278
Epoch [26], train_loss: 0.9874, val_loss: 1.0638, val_acc: 0.6238
Epoch [27], train_loss: 0.9749, val_loss: 1.0492, val_acc: 0.6372
Epoch [28], train_loss: 0.9664, val_loss: 1.0464, val_acc: 0.6354
Epoch [29], train_loss: 0.9569, val_loss: 1.0394, val_acc: 0.6360
Epoch [30], train_loss: 0.9489, val_loss: 1.0308, val_acc: 0.6392
Epoch [31], train_loss: 0.9408, val_loss: 1.0229, val_acc: 0.6440
Epoch [32], train_loss: 0.9332, val_loss: 1.0293, val_acc: 0.6275
Epoch [33], train_loss: 0.9194, val_loss: 1.0411, val_acc: 0.6502
Epoch [34], train_loss: 0.9188, val_loss: 1.0381, val_acc: 0.6419
Epoch [35], train_loss: 0.9046, val_loss: 1.0220, val_acc: 0.6591
Epoch [36], train_loss: 0.8986, val_loss: 1.0075, val_acc: 0.6462
Epoch [37], train_loss: 0.8956, val_loss: 1.0094, val_acc: 0.6508
Epoch [38], train_loss: 0.8789, val_loss: 0.9879, val_acc: 0.6592
Epoch [39], train_loss: 0.8811, val_loss: 0.9358, val_acc: 0.6517
Epoch [40], train_loss: 0.8723, val_loss: 1.0092, val_acc: 0.6590
Epoch [41], train_loss: 0.8589, val_loss: 0.9575, val_acc: 0.6701
Epoch [42], train_loss: 0.8575, val_loss: 1.0032, val_acc: 0.6530
Epoch [43], train_loss: 0.8521, val_loss: 0.9722, val_acc: 0.6688
Epoch [44], train_loss: 0.8458, val_loss: 0.9909, val_acc: 0.6677
Epoch [45], train_loss: 0.8346, val_loss: 0.9617, val_acc: 0.6694
Epoch [46], train_loss: 0.8303, val_loss: 0.9926, val_acc: 0.6573
Epoch [47], train_loss: 0.8267, val_loss: 0.9877, val_acc: 0.6672
Epoch [48], train_loss: 0.8192, val_loss: 0.9642, val_acc: 0.6675
Epoch [49], train_loss: 0.8249, val_loss: 0.9295, val_acc: 0.6682
Epoch [50], train_loss: 0.7962, val_loss: 0.9696, val_acc: 0.6644
Epoch [51], train_loss: 0.7917, val_loss: 0.9471, val_acc: 0.6768
Epoch [52], train_loss: 0.7855, val_loss: 0.9295, val_acc: 0.6692
Epoch [53], train_loss: 0.7780, val_loss: 0.9442, val_acc: 0.6705
Epoch [54], train_loss: 0.7713, val_loss: 0.9770, val_acc: 0.6591
Epoch [55], train_loss: 0.7652, val_loss: 0.9394, val_acc: 0.6734
Epoch [56], train_loss: 0.7689, val_loss: 0.9411, val_acc: 0.6702
Epoch [57], train_loss: 0.7524, val_loss: 0.9568, val_acc: 0.6704
Epoch [58], train_loss: 0.7483, val_loss: 0.9504, val_acc: 0.6699
Epoch [59], train_loss: 0.7422, val_loss: 0.9398, val_acc: 0.6707
Epoch [60], train_loss: 0.7352, val_loss: 0.9255, val_acc: 0.6873
Epoch [61], train_loss: 0.7290, val_loss: 0.9225, val_acc: 0.6824
Epoch [62], train_loss: 0.7189, val_loss: 0.9269, val_acc: 0.6802
Epoch [63], train_loss: 0.7169, val_loss: 0.9252, val_acc: 0.6843
Epoch [64], train_loss: 0.7117, val_loss: 0.9379, val_acc: 0.6841
Epoch [65], train_loss: 0.7053, val_loss: 0.9434, val_acc: 0.6707
Epoch [66], train_loss: 0.6926, val_loss: 0.9331, val_acc: 0.6928
Epoch [67], train_loss: 0.6909, val_loss: 0.9325, val_acc: 0.6911
Epoch [68], train_loss: 0.6837, val_loss: 0.9332, val_acc: 0.6915
Epoch [69], train_loss: 0.6786, val_loss: 0.9146, val_acc: 0.6936
Epoch [70], train_loss: 0.6744, val_loss: 0.9255, val_acc: 0.6902
Epoch [71], train_loss: 0.6747, val_loss: 0.9189, val_acc: 0.6905
Epoch [72], train_loss: 0.6654, val_loss: 0.9223, val_acc: 0.6939
Epoch [73], train_loss: 0.6654, val_loss: 0.9263, val_acc: 0.6930
Epoch [74], train_loss: 0.6587, val_loss: 0.9340, val_acc: 0.6910
Epoch [75], train_loss: 0.6462, val_loss: 0.9354, val_acc: 0.6934
Epoch [76], train_loss: 0.6396, val_loss: 0.9316, val_acc: 0.6914
Epoch [77], train_loss: 0.6342, val_loss: 0.9266, val_acc: 0.6934
Epoch [78], train_loss: 0.6294, val_loss: 0.9485, val_acc: 0.6969
Epoch [79], train_loss: 0.6281, val_loss: 0.9340, val_acc: 0.6991
Epoch [80], train_loss: 0.6260, val_loss: 0.9352, val_acc: 0.6996
Epoch [81], train_loss: 0.6195, val_loss: 0.9079, val_acc: 0.6991
Epoch [82], train_loss: 0.6074, val_loss: 0.9133, val_acc: 0.6964
Epoch [83], train_loss: 0.6087, val_loss: 0.9127, val_acc: 0.7007
Epoch [84], train_loss: 0.6081, val_loss: 0.9146, val_acc: 0.6914
Epoch [85], train_loss: 0.5990, val_loss: 0.9389, val_acc: 0.6914
Epoch [86], train_loss: 0.5928, val_loss: 0.9139, val_acc: 0.7013
Epoch [87], train_loss: 0.5862, val_loss: 0.9272, val_acc: 0.6987
Epoch [88], train_loss: 0.5855, val_loss: 0.9190, val_acc: 0.6998
Epoch [89], train_loss: 0.5793, val_loss: 0.9205, val_acc: 0.6947
Epoch [90], train_loss: 0.5785, val_loss: 0.9268, val_acc: 0.7010
Epoch [91], train_loss: 0.5696, val_loss: 0.9288, val_acc: 0.6982
Epoch [92], train_loss: 0.5654, val_loss: 0.9358, val_acc: 0.6948
Epoch [93], train_loss: 0.5592, val_loss: 0.9248, val_acc: 0.6965
Epoch [94], train_loss: 0.5494, val_loss: 0.9273, val_acc: 0.6986
Epoch [95], train_loss: 0.5444, val_loss: 0.9354, val_acc: 0.6934
Epoch [96], train_loss: 0.5396, val_loss: 0.9316, val_acc: 0.6914
Epoch [97], train_loss: 0.5342, val_loss: 0.9266, val_acc: 0.6934
Epoch [98], train_loss: 0.5294, val_loss: 0.9485, val_acc: 0.6969
Epoch [99], train_loss: 0.5281, val_loss: 0.9340, val_acc: 0.6991
Epoch [100], train_loss: 0.5225, val_loss: 0.9305, val_acc: 0.6995
Epoch [101], train_loss: 0.5131, val_loss: 0.9407, val_acc: 0.7028
Epoch [102], train_loss: 0.5089, val_loss: 0.9272, val_acc: 0.6992
Epoch [103], train_loss: 0.5046, val_loss: 0.9318, val_acc: 0.6991
Epoch [104], train_loss: 0.5009, val_loss: 0.9401, val_acc: 0.6982
Epoch [105], train_loss: 0.4925, val_loss: 0.9050, val_acc: 0.6992
Epoch [106], train_loss: 0.4846, val_loss: 0.9409, val_acc: 0.6960
Epoch [107], train_loss: 0.4783, val_loss: 0.9432, val_acc: 0.6914
Epoch [108], train_loss: 0.4762, val_loss: 0.9432, val_acc: 0.7014
Epoch [109], train_loss: 0.4715, val_loss: 0.9482, val_acc: 0.7005
Epoch [110], train_loss: 0.4681, val_loss: 0.9502, val_acc: 0.7023
Epoch [111], train_loss: 0.4639, val_loss: 0.9516, val_acc: 0.6985
Epoch [112], train_loss: 0.4559, val_loss: 0.9710, val_acc: 0.6920
Epoch [113], train_loss: 0.4517, val_loss: 0.9619, val_acc: 0.7008
Epoch [114], train_loss: 0.4456, val_loss: 0.9608, val_acc: 0.7054
Epoch [115], train_loss: 0.4460, val_loss: 0.9628, val_acc: 0.6965
Epoch [116], train_loss: 0.4373, val_loss: 0.9884, val_acc: 0.7015
Epoch [117], train_loss: 0.4340, val_loss: 0.9735, val_acc: 0.6991
Epoch [118], train_loss: 0.4276, val_loss: 0.9632, val_acc: 0.7011
Epoch [119], train_loss: 0.4230, val_loss: 0.9920, val_acc: 0.6968
Epoch [120], train_loss: 0.4192, val_loss: 0.9712, val_acc: 0.6972
Epoch [121], train_loss: 0.4146, val_loss: 0.9883, val_acc: 0.6997
Epoch [122], train_loss: 0.4111, val_loss: 1.0082, val_acc: 0.6963
Epoch [123], train_loss: 0.4086, val_loss: 0.9759, val_acc: 0.6992
Epoch [124], train_loss: 0.4026, val_loss: 0.9715, val_acc: 0.6997
Epoch [125], train_loss: 0.3955, val_loss: 0.9793, val_acc: 0.7034
Epoch [126], train_loss: 0.3935, val_loss: 1.0134, val_acc: 0.6997
Epoch [127], train_loss: 0.3889, val_loss: 0.9819, val_acc: 0.6992
Epoch [128], train_loss: 0.3815, val_loss: 1.0071, val_acc: 0.6986
Epoch [129], train_loss: 0.3774, val_loss: 1.0262, val_acc: 0.6986
Epoch [130], train_loss: 0.3740, val_loss: 1.0409, val_acc: 0.6986
Epoch [131], train_loss: 0.3639, val_loss: 1.0225, val_acc: 0.7045
Epoch [132], train_loss: 0.3675, val_loss: 1.0216, val_acc: 0.7045
Epoch [133], train_loss: 0.3545, val_loss: 1.0523, val_acc: 0.6997
Epoch [134], train_loss: 0.3542, val_loss: 1.0394, val_acc: 0.7008
Epoch [135], train_loss: 0.3488, val_loss: 1.0577, val_acc: 0.6977
Epoch [136], train_loss: 0.3422, val_loss: 1.0812, val_acc: 0.6900
Epoch [137], train_loss: 0.3395, val_loss: 1.0594, val_acc: 0.6944
Epoch [138], train_loss: 0.3349, val_loss: 1.0533, val_acc: 0.6962
Epoch [139], train_loss: 0.3291, val_loss: 1.0960, val_acc: 0.6944
Epoch [140], train_loss: 0.3274, val_loss: 1.0971, val_acc: 0.7034
Epoch [141], train_loss: 0.3242, val_loss: 1.0799, val_acc: 0.6973
Epoch [142], train_loss: 0.3162, val_loss: 1.0739, val_acc: 0.6973
Epoch [143], train_loss: 0.3181, val_loss: 1.0787, val_acc: 0.7041
Epoch [144], train_loss: 0.3089, val_loss: 1.1401, val_acc: 0.6982
Epoch [145], train_loss: 0.3011, val_loss: 1.1108, val_acc: 0.6973
Epoch [146], train_loss: 0.2971, val_loss: 1.0753, val_acc: 0.7016
Epoch [147], train_loss: 0.2889, val_loss: 1.1518, val_acc: 0.6977
Epoch [148], train_loss: 0.2858, val_loss: 1.1219, val_acc: 0.6974
Epoch [149], train_loss: 0.2850, val_loss: 1.0999, val_acc: 0.7070
Epoch [150], train_loss: 0.2821, val_loss: 1.1212, val_acc: 0.6991
Epoch [151], train_loss: 0.2777, val_loss: 1.1313, val_acc: 0.6994
Epoch [152], train_loss: 0.2725, val_loss: 1.1371, val_acc: 0.6984
Epoch [153], train_loss: 0.2692, val_loss: 1.1482, val_acc: 0.6977
Epoch [154], train_loss: 0.2662, val_loss: 1.1479, val_acc: 0.6976
Epoch [155], train_loss: 0.2614, val_loss: 1.1341, val_acc: 0.6976
Epoch [156], train_loss: 0.2592, val_loss: 1.1259, val_acc: 0.6976
Epoch [157], train_loss: 0.2533, val_loss: 1.1585, val_acc: 0.6992
Epoch [158], train_loss: 0.2533, val_loss: 1.1557, val_acc: 0.6974
Epoch [159], train_loss: 0.2459, val_loss: 1.1576, val_acc: 0.6984
Epoch [160], train_loss: 0.2409, val_loss: 1.1482, val_acc: 0.6982
Epoch [161], train_loss: 0.2396, val_loss: 1.1664, val_acc: 0.6996
Epoch [162], train_loss: 0.2350, val_loss: 1.1692, val_acc: 0.6972
Epoch [163], train_loss: 0.2349, val_loss: 1.1764, val_acc: 0.6982
Epoch [164], train_loss: 0.2303, val_loss: 1.1852, val_acc: 0.7020
Epoch [165], train_loss: 0.2231, val_loss: 1.2016, val_acc: 0.6965
Epoch [166], train_loss: 0.2206, val_loss: 1.2235, val_acc: 0.6965
Epoch [167], train_loss: 0.2132, val_loss: 1.2373, val_acc: 0.6964
Epoch [168], train_loss: 0.2091, val_loss: 1.2454, val_acc: 0.6993
Epoch [169], train_loss: 0.2056, val_loss: 1.2565, val_acc: 0.6993
Epoch [170], train_loss: 0.2009, val_loss: 1.2575, val_acc: 0.6999
Epoch [171], train_loss: 0.1955, val_loss: 1.2712, val_acc: 0.6987
Epoch [172], train_loss: 0.1895, val_loss: 1.2697, val_acc: 0.6983
Epoch [173], train_loss: 0.1800, val_loss: 1.2903, val_acc: 0.6933
Epoch [174], train_loss: 0.1809, val_loss: 1.2985, val_acc: 0.6947
Epoch [175], train_loss: 0.1789, val_loss: 1.3437, val_acc: 0.6982
Epoch [176], train_loss: 0.1729, val_loss: 1.3036, val_acc: 0.6972
Epoch [177], train_loss: 0.1733, val_loss: 1.3229, val_acc: 0.6976
Epoch [178], train_loss: 0.1688, val_loss: 1.2985, val_acc: 0.6986
Epoch [179], train_loss: 0.1644, val_loss: 1.3260, val_acc: 0.6986
Epoch [180], train_loss: 0.1601, val_loss: 1.3351, val_acc: 0.7009
Epoch [181], train_loss: 0.1588, val_loss: 1.3178, val_acc: 0.6991
Epoch [182], train_loss: 0.1539, val_loss: 1.3436, val_acc: 0.6969
Epoch [183], train_loss: 0.1533, val_loss: 1.3473, val_acc: 0.6984
Epoch [184], train_loss: 0.1464, val_loss: 1.3609, val_acc: 0.6982
Epoch [185], train_loss: 0.1452, val_loss: 1.3521, val_acc: 0.6950
Epoch [186], train_loss: 0.1407, val_loss: 1.3764, val_acc: 0.6977
Epoch [187], train_loss: 0.1372, val_loss: 1.3972, val_acc: 0.6975
Epoch [188], train_loss: 0.1363, val_loss: 1.3924, val_acc: 0.6975
Epoch [189], train_loss: 0.1348, val_loss: 1.4206, val_acc: 0.6985
Epoch [190], train_loss: 0.1279, val_loss: 1.4106, val_acc: 0.6979
```

```
In [2]: def predict_image(img, model):
    # Convert to batch of 1
    xb = to_device(img.unsqueeze(0), device)
    # Get prediction from model
    yb = model(xb)
    # Pick index with highest probability
    _, preds = torch.max(yb, dim=1)
    # Retrieve the class label
    return dataset.classes[preds[0].item()]

In [43]: img, label = test_dataset[0]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))

Label: airplane, Predicted: airplane

0
5
10
15
20
25
30
0 5 10 15 20 25 30

In [44]: img, label = test_dataset[1002]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))

Label: automobile, Predicted: truck

0
5
10
15
20
25
30
0 5 10 15 20 25 30
```

이러한 결과는 모델이 자동차를 잘못 분류한 것 같습니다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.

## Testing with individual images

```
In [41]: test_dataset = ImageFolder(data_dir+'/'+'test', transform=ToTensor())

In [42]: def predict_image(img, model):
    # Convert to batch of 1
    xb = to_device(img.unsqueeze(0), device)
    # Get prediction from model
    yb = model(xb)
    # Pick index with highest probability
    _, preds = torch.max(yb, dim=1)
    # Retrieve the class label
    return dataset.classes[preds[0].item()]

In [43]: img, label = test_dataset[0]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))

Label: airplane, Predicted: airplane

0
5
10
15
20
25
30
0 5 10 15 20 25 30

In [44]: img, label = test_dataset[1002]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))

Label: automobile, Predicted: truck

0
5
10
15
20
25
30
0 5 10 15 20 25 30
```

이러한 결과는 모델이 자동차를 잘못 분류한 것 같습니다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.

이 과정의 결과를 위해 "노이즈"를 추가한다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.

## Testing with individual images

```
In [41]: test_dataset = ImageFolder(data_dir+'/'+'test', transform=ToTensor())

In [42]: def predict_image(img, model):
    # Convert to batch of 1
    xb = to_device(img.unsqueeze(0), device)
    # Get prediction from model
    yb = model(xb)
    # Pick index with highest probability
    _, preds = torch.max(yb, dim=1)
    # Retrieve the class label
    return dataset.classes[preds[0].item()]

In [43]: img, label = test_dataset[0]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))

Label: airplane, Predicted: airplane

0
5
10
15
20
25
30
0 5 10 15 20 25 30

In [44]: img, label = test_dataset[1002]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))

Label: automobile, Predicted: truck

0
5
10
15
20
25
30
0 5 10 15 20 25 30
```

이러한 결과는 모델이 자동차를 잘못 분류한 것 같습니다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.

이 과정의 결과를 위해 "노이즈"를 추가한다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.

이 과정의 결과를 위해 "노이즈"를 추가한다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.

이 과정의 결과를 위해 "노이즈"를 추가한다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.

이 과정의 결과를 위해 "노이즈"를 추가한다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.

이 과정의 결과를 위해 "노이즈"를 추가한다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해