

Cifar10-CNN ipynb version

Image Classification using Convolutional Neural Networks(CNN) in PyTorch

1. Exploring the Cifar10 Dataset

```
In [ ]: import os
import torch
import torchvision
import tarfile
from torchvision.datasets.utils import download_url
from torch.utils.data import random_split
```

```
In [110... # Download the dataset
dataset_url = "https://s3.amazonaws.com/fast-ai-imageclas/cifar10.tgz"
download_url(dataset_url, '.')
```

Using downloaded and verified file: .\cifar10.tgz

```
In [111... # Extract from archive
with tarfile.open('./cifar10.tgz', 'r:gz') as tar:
    tar.extractall(path='./data')
```

```
In [112... data_dir = './data/cifar10'
```

```
print(os.listdir(data_dir))
classes = os.listdir(data_dir + "/train")
print(classes)
```

```
['test', 'train']
['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

두개의 폴더 내부를 살펴보면, 하나는 train, test이고 각 클래스에 동일한 수의 이미지가 있는지 확인. train은 5000개, test는 1000개

```
In [113... airplane_files = os.listdir(data_dir + "/train/airplane")
print('No. of training examples for airplanes:', len(airplane_files))
print(airplane_files[:5])
```

```
No. of training examples for airplanes: 5000
['0001.png', '0002.png', '0003.png', '0004.png', '0005.png']
```

```
In [114... ship_test_files = os.listdir(data_dir + "/test/ship")
print("No. of test examples for ship:", len(ship_test_files))
print(ship_test_files[:5])
```

```
No. of test examples for ship: 1000
['0001.png', '0002.png', '0003.png', '0004.png', '0005.png']
```

이미지 폴더의 class를 torchvision을 통해 pytorch tensor로 로드해본다.

```
In [115... from torchvision.datasets import ImageFolder
from torchvision.transforms import ToTensor
```

```
In [116... dataset = ImageFolder(data_dir+'train', transform = ToTensor())
```

각 원소는 튜플이고 이미지 텐서와 라벨을 갖고있다.

cf. 튜플 : 순서가 있는 객체의 집합으로 리스트와 유사하다. 하지만 값 변경 X

- 이미지 텐서는 32*32 pixel 에 channel은 3(RGB)라서 각 이미지 shape = (3,32,32)

```
In [117... img, label = dataset[0]
print(img.shape, label)
img
```

```
torch.Size([3, 32, 32]) 0
```

```
Out[117]: tensor([[0.7922, 0.7922, 0.8000, ..., 0.8118, 0.8039, 0.7961],
        [0.8078, 0.8078, 0.8118, ..., 0.8235, 0.8157, 0.8078],
        [0.8235, 0.8275, 0.8314, ..., 0.8392, 0.8314, 0.8235],
        ...,
        [0.8549, 0.8235, 0.7608, ..., 0.9529, 0.9569, 0.9529],
        [0.8588, 0.8510, 0.8471, ..., 0.9451, 0.9451, 0.9451],
        [0.8510, 0.8471, 0.8510, ..., 0.9373, 0.9373, 0.9412]],

        [[0.8000, 0.8000, 0.8078, ..., 0.8157, 0.8078, 0.8000],
        [0.8157, 0.8157, 0.8196, ..., 0.8275, 0.8196, 0.8118],
        [0.8314, 0.8353, 0.8392, ..., 0.8392, 0.8353, 0.8275],
        ...,
        [0.8510, 0.8196, 0.7608, ..., 0.9490, 0.9490, 0.9529],
        [0.8549, 0.8471, 0.8471, ..., 0.9412, 0.9412, 0.9412],
        [0.8471, 0.8431, 0.8471, ..., 0.9333, 0.9333, 0.9333]],

        [[0.7804, 0.7804, 0.7882, ..., 0.7843, 0.7804, 0.7765],
        [0.7961, 0.7961, 0.8000, ..., 0.8039, 0.7961, 0.7882],
        [0.8118, 0.8157, 0.8235, ..., 0.8235, 0.8157, 0.8078],
        ...,
        [0.8706, 0.8392, 0.7765, ..., 0.9686, 0.9686, 0.9686],
        [0.8745, 0.8667, 0.8627, ..., 0.9608, 0.9608, 0.9608],
        [0.8667, 0.8627, 0.8667, ..., 0.9529, 0.9529, 0.9529]])
```

```
In [118]: print(dataset.classes)
```

```
['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

여기 ipynb 상에서 시각화를 하고싶다면 matplotlib을 써야하는데 이때는 텐서 차원을 (32x32x3)으로 바꿔야함


```
In [119]: import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
# 요건 그림 위에 수식 없이 바로 보여지게끔 하는 코딩
matplotlib.rcParams['figure.facecolor'] = '#ffffff'
```

```
In [120]: def show_example(img, label):
    print('Label: ', dataset.classes[label], "(" + str(label) + ")")
    plt.imshow(img.permute(1,2,0))

# """
# # Permute() : 모든 차원들을 맞교환할 수 있음.
# ex)
# x = torch.rand(16, 32, 3)
# y = x.permute(2,1,0)    --> [3, 32, 16]
# """
```


```
In [121]: show_example(*dataset[0])
```

```
Label: airplane (0)
```



```
In [122]: show_example(*dataset[10000])
```

```
Label: bird (2)
```



Train, Valid Split

```
In [123.. random_seed = 42
torch.manual_seed(random_seed);
```

```
In [124.. val_size = 5000
train_size = len(dataset) - val_size

train_ds, val_ds = random_split(dataset, [train_size, val_size])
len(train_ds), len(val_ds)
```

Out[124]: (45000, 5000)

```
In [125.. from torch.utils.data.dataloader import DataLoader

batch_size=512
```


배치사이즈에 load 시켜주려면 daata loader가 필요하니까

```
In [126.. train_dl = DataLoader(train_ds, batch_size, shuffle=True, num_workers=4, pin_memory=True)
val_dl = DataLoader(val_ds, batch_size*2, num_workers=4, pin_memory=True)
```

```
In [127.. from torchvision.utils import make_grid

def show_batch(dl):
    for images, labels in dl:
        fig, ax = plt.subplots(figsize=(12,6))
        ax.set_xticks([]); ax.set_yticks([])
        ax.imshow(make_grid(images, nrow=16).permute(1, 2, 0))
        break
```

```
In [128.. show_batch(train_dl)
```



Defining the Model(Convolutional Neural Network)

```
In [129.. def apply_kernel(image, kernel):
    ri, ci = image.shape           # image dimensions
    rk, ck = kernel.shape         # kernel dimensions
    ro, co = ri-rk+1, ci-ck+1    # output dimensions
    output = torch.zeros([ro, co])
    for i in range(ro):
        for j in range(co):
            output[i,j] = torch.sum(image[i:i+rk, j:j+ck] * kernel)

    return output
```

```
In [130.. import torch.nn as nn
import torch.nn.functional as F
```

```
In [131.. simple_model = nn.Sequential(
    nn.Conv2d(3, 8, kernel_size=3, stride=1, padding=1),
    nn.MaxPool2d(2, 2)
)
```

```
In [132.. for images, labels in train_dl:
    print('images.shape:', images.shape)
    out = simple_model(images)
    print('out.shape:', out.shape)
    break
```

images.shape: torch.Size([512, 3, 32, 32])
out.shape: torch.Size([512, 8, 16, 16])



```
In [133.. class ImageClassificationBase(nn.Module):
    def training_step(self, batch):
        images, labels = batch
        out = self(images)           # Generate predictions
        loss = F.cross_entropy(out, labels) # Calculate loss
        return loss

    def validation_step(self, batch):
        images, labels = batch
        out = self(images)           # Generate predictions
        loss = F.cross_entropy(out, labels) # Calculate loss
        acc = accuracy(out, labels)     # Calculate accuracy
        return {'val_loss': loss.detach(), 'val_acc': acc}

    def validation_epoch_end(self, outputs):
        batch_losses = [x['val_loss'] for x in outputs]
        epoch_loss = torch.stack(batch_losses).mean() # Combine losses
        batch_accs = [x['val_acc'] for x in outputs]
        epoch_acc = torch.stack(batch_accs).mean() # Combine accuracies
        return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}

    def epoch_end(self, epoch, result):
```

```

        print("Epoch [{}], train_loss: {:.4f}, val_loss: {:.4f}, val_acc: {:.4f}".format(
            epoch, result['train_loss'], result['val_loss'], result['val_acc']))

def accuracy(outputs, labels):
    _, preds = torch.max(outputs, dim=1)
    return torch.tensor(torch.sum(preds == labels).item() / len(preds))

```

nn.sequential 을 사용하여 층과 활성화함수를 하나의 network architecture로 연결시킬것이다.

```

In [134]: class Cifar10CnnModel(ImageClassificationBase):
            def __init__(self):
                super().__init__()
                self.network = nn.Sequential(
                    nn.Conv2d(3, 32, kernel_size=3, padding=1),
                    nn.ReLU(),
                    nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
                    nn.ReLU(),
                    nn.MaxPool2d(2, 2), # output: 64 x 16 x 16

                    nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
                    nn.ReLU(),
                    nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1),
                    nn.ReLU(),
                    nn.MaxPool2d(2, 2), # output: 128 x 8 x 8

                    nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1),
                    nn.ReLU(),
                    nn.Conv2d(256, 256, kernel_size=3, stride=1, padding=1),
                    nn.ReLU(),
                    nn.MaxPool2d(2, 2), # output: 256 x 4 x 4

                    nn.Flatten(),
                    nn.Linear(256*4*4, 1024),
                    nn.ReLU(),
                    nn.Linear(1024, 512),
                    nn.ReLU(),
                    nn.Linear(512, 10))

            def forward(self, xb):
                return self.network(xb)

```

```

In [135]: model = Cifar10CnnModel()
            model

```

```

Out[135]: Cifar10CnnModel(
  (network): Sequential(
    (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU()
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU()
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU()
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU()
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU()
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (15): Flatten(start_dim=1, end_dim=-1)
    (16): Linear(in_features=4096, out_features=1024, bias=True)
    (17): ReLU()
    (18): Linear(in_features=1024, out_features=512, bias=True)
    (19): ReLU()
    (20): Linear(in_features=512, out_features=10, bias=True)
  )
)

```

```

In [136]: for images, labels in train_dl:
            print('images.shape:', images.shape)
            out = model(images)
            print('out.shape:', out.shape)
            print('out[0]:', out[0])
            break

```

```

images.shape: torch.Size([512, 3, 32, 32])
out.shape: torch.Size([512, 10])
out[0]: tensor([ 0.0239, -0.0466,  0.0067,  0.0193,  0.0044, -0.0598, -0.0188, -0.0242,
                  0.0431, -0.0164], grad_fn=<SelectBackward0>)

```

GPU를 원활하게 사용하기 위해 사용 가능한 경우 몇 가지 도우미 함수(get_default_device 및 to_device)와 도우미 클래스 DeviceDataLoader를 정의하여 필요에 따라 모델 및 데이터를 GPU로 이동시킨다!

```

In [137]: def get_default_device():
            """Pick GPU if available, else CPU"""

```

```

    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    """Move tensor(s) to chosen device"""
    if isinstance(data, (list,tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    """Wrap a dataloader to move data to a device"""
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        """Yield a batch of data after moving it to device"""
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        """Number of batches"""
        return len(self.dl)

```

```

In [138]: device = get_default_device()
          device

```

```

Out[138]: device(type='cuda')

```

이제 DeviceDataLoader를 사용하여 데이터 배치를 GPU로 자동 전송(사용 가능한 경우)하고 to_device를 사용하여 모델을 GPU(사용 가능한 경우)로 이동하기 위해 교육 및 검증 데이터 로더를 래핑할 수 있습니다.

```

In [139]: train_dl = DeviceDataLoader(train_dl, device)
          val_dl = DeviceDataLoader(val_dl, device)
          to_device(model, device);

```

Training the Model

```

In [140]: @torch.no_grad()
          def evaluate(model, val_loader):
              model.eval()
              outputs = [model.validation_step(batch) for batch in val_loader]
              return model.validation_epoch_end(outputs)

          def fit(epochs, lr, model, train_loader, val_loader, opt_func=torch.optim.SGD):
              history = []
              optimizer = opt_func(model.parameters(), lr)
              for epoch in range(epochs):
                  # Training Phase
                  model.train()
                  train_losses = []
                  for batch in train_loader:
                      loss = model.training_step(batch)
                      train_losses.append(loss)
                      loss.backward()
                      optimizer.step()
                      optimizer.zero_grad()
                  # Validation phase
                  result = evaluate(model, val_loader)
                  result['train_loss'] = torch.stack(train_losses).mean().item()
                  model.epoch_end(epoch, result)
                  history.append(result)
              return history

```

```

In [141]: model = to_device(Cifar10CnnModel(), device)

```

```

In [142]: evaluate(model, val_dl)

```

```

Out[142]: {'val_loss': 2.3023083209991455, 'val_acc': 0.09984616935253143}

```

무작위에서 뽑았을때(초기정확도)는 10% 이며, 우리는 모델을 훈련하기 위해 하이퍼파라미터(학습속도, epoch수, batch_size 등)를 사용할 것이다. -> 높은 정확도를 위하여

```

In [155]: num_epochs = 30
          opt_func = torch.optim.Adam
          lr = 0.001

```

```


In [156]: history = fit(num_epochs, lr, model, train_dl, val_dl, opt_func)

```

```
Epoch [0], train_loss: 0.7214, val_loss: 0.7417, val_acc: 0.7558
Epoch [1], train_loss: 0.4055, val_loss: 0.6987, val_acc: 0.7728
Epoch [2], train_loss: 0.3318, val_loss: 0.7446, val_acc: 0.7632
Epoch [3], train_loss: 0.2621, val_loss: 0.7991, val_acc: 0.7642
Epoch [4], train_loss: 0.2087, val_loss: 0.8144, val_acc: 0.7734
Epoch [5], train_loss: 0.1569, val_loss: 0.9376, val_acc: 0.7721
Epoch [6], train_loss: 0.1147, val_loss: 1.0520, val_acc: 0.7643
Epoch [7], train_loss: 0.0973, val_loss: 1.0965, val_acc: 0.7660
Epoch [8], train_loss: 0.0700, val_loss: 1.2361, val_acc: 0.7685
Epoch [9], train_loss: 0.0696, val_loss: 1.1679, val_acc: 0.7719
Epoch [10], train_loss: 0.0564, val_loss: 1.4161, val_acc: 0.7569
Epoch [11], train_loss: 0.0624, val_loss: 1.2999, val_acc: 0.7677
Epoch [12], train_loss: 0.0468, val_loss: 1.4434, val_acc: 0.7491
Epoch [13], train_loss: 0.0427, val_loss: 1.4172, val_acc: 0.7684
Epoch [14], train_loss: 0.0405, val_loss: 1.3459, val_acc: 0.7712
Epoch [15], train_loss: 0.0451, val_loss: 1.5065, val_acc: 0.7667
Epoch [16], train_loss: 0.0383, val_loss: 1.3871, val_acc: 0.7707
Epoch [17], train_loss: 0.0315, val_loss: 1.5069, val_acc: 0.7676
Epoch [18], train_loss: 0.0370, val_loss: 1.4604, val_acc: 0.7760
Epoch [19], train_loss: 0.0396, val_loss: 1.3977, val_acc: 0.7698
Epoch [20], train_loss: 0.0288, val_loss: 1.5666, val_acc: 0.7640
Epoch [21], train_loss: 0.0337, val_loss: 1.5082, val_acc: 0.7697
Epoch [22], train_loss: 0.0331, val_loss: 1.5070, val_acc: 0.7731
Epoch [23], train_loss: 0.0252, val_loss: 1.6157, val_acc: 0.7688
Epoch [24], train_loss: 0.0499, val_loss: 1.5263, val_acc: 0.7643
Epoch [25], train_loss: 0.0206, val_loss: 1.5882, val_acc: 0.7804
Epoch [26], train_loss: 0.0371, val_loss: 1.4509, val_acc: 0.7730
Epoch [27], train_loss: 0.0290, val_loss: 1.5327, val_acc: 0.7702
Epoch [28], train_loss: 0.0194, val_loss: 1.7540, val_acc: 0.7740
Epoch [29], train_loss: 0.0333, val_loss: 1.5058, val_acc: 0.7739
```

```
In [157.. def plot_accuracies(history):
    accuracies = [x['val_acc'] for x in history]
    plt.plot(accuracies, '-x')
    plt.xlabel('epoch')
    plt.ylabel('accuracy')
    plt.title('Accuracy vs. No. of epochs');
```


```
In [158.. plot_accuracies(history)
```



이 모델은 약 75% 이상 정도의 정확도에 도달하였고, epoch을 더 돌린다 하더라도 정확도가 올라갈 것 같지 않다.

```
In [159.. def plot_losses(history):
    train_losses = [x.get('train_loss') for x in history]
    val_losses = [x['val_loss'] for x in history]
    plt.plot(train_losses, '-bx')
    plt.plot(val_losses, '-rx')
    plt.xlabel('epoch')
    plt.ylabel('loss')
    plt.legend(['Training', 'Validation'])
    plt.title('Loss vs. No. of epochs');
```

```
In [160.. plot_losses(history)
```



처음에는 훈련과 검증 모두 손실이 낮아지다가 epoch 3부터는 훈련손실만 더 낮아지고 검증set 손실은 더 높아지는것을 확인할 수 있다. - > 과적합!

이 과적합을 피하기 위해 "노이즈"를 추가한다. 배치 정규화 및 드롭아웃과 같은 정규화 기술 사용을 통해.


Testing with individual images

```
In [149.. test_dataset = ImageFolder(data_dir+'/test', transform=ToTensor())
```

```
In [150.. def predict_image(img, model):
# Convert to a batch of 1
xb = to_device(img.unsqueeze(0), device)
# Get predictions from model
yb = model(xb)
# Pick index with highest probability
_, preds = torch.max(yb, dim=1)
# Retrieve the class label
return dataset.classes[preds[0].item()]
```


```
In [151.. img, label = test_dataset[0]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))
```

Label: airplane , Predicted: airplane




```
In [152.. img, label = test_dataset[1002]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))
```

Label:automobile , Predicted: automobile



```
In [153]: img, label = test_dataset[6153]
plt.imshow(img.permute(1, 2, 0))
print('Label:', dataset.classes[label], ', Predicted:', predict_image(img, model))
```

Label: frog , Predicted: frog



```
In [154]: test_loader = DeviceDataLoader(DataLoader(test_dataset, batch_size*2), device)
result = evaluate(model, test_loader)
result
```

```
Out[154]: {'val_loss': 0.7234663963317871, 'val_acc': 0.7527562975883484}
```

Saving and loading the model

```
In [161]: torch.save(model.state_dict(), 'cifar10-cnn-acc0.7527-epoch10-lr0.001-adam-bs512.pth')
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js